


Introducing Classes

A **class** is an **interface** paired with an **implementation**, similar to the **Time** structure you created in the lesson on Information Hiding. The **public interface** specifies how clients interact with objects, and the **private implementation** specifies how the functions in the public interface are implemented.



The **Time** struct (even when paired with an interface, so the data is hidden), still allows users to **directly manipulate** its data members. Classes take a different approach; with classes, the data inside your objects will **only be accessible by the member functions**, **forcing** the client to access and modify data in a safe way.

A Class Definition

Here is a **class definition** similar to the structure from last week (and **H27**):

```
class Time
{
public:
    int getHours() const;
    int getMinutes() const;
    Time sum(const Time& rhs) const;
    Time difference(const Time& rhs) const;

    istream& read(istream& in);
    ostream& print(ostream& out) const;
private:
    int m_hours;
    int m_minutes;
};
```

1. Instead of **struct**, use the **class** keyword. There is no **public** in front.
2. The **public** keyword, followed by a colon, indicates the start of the **public interface**. Here we **prototype** the member functions it contains.
3. The member functions **getHours()**, **getMinutes()**, **sum()**, **difference()** and **print()**, all access the **hours** and **minutes** **without changing them**. Add the **const** keyword after the argument list. We say these are **accessors**.
4. The **read()** member function **modifies** the **Time** object. This is a **mutator**.
5. The class definition **ends with a semicolon**. This is not optional.

Data Members



Most of the implementation will appear inside a **.cpp** file. Defining the data members which store **object state**, is **written inside the header file instead**. A common practice is to use a special indicator like **m_** to show that it is a data member.

The **Time struct** used two individual data members: one for hours and one for minutes. This is fine; it allows you to store all of information needed. By adding **private**, you can **prevent clients of Time from accessing the fields directly**.

Introducing Encapsulation

So, what do **public** and **private** mean? If a member of a class is **public**, then **any part of your code** can access and manipulate it. If you have a **public** member function, any code can call it **using an object of that type**. If a data member is marked **private**, then only member functions of the class can access it.

The **public** and **private** keywords are the C++ mechanism for **defining interfaces and enforcing encapsulation**. Once you add **private**, the compiler enforces the **appropriate encapsulation**.

```

5
6 int main()
7 {
8     Time t;           // a Time object
9     t.hours = 18;     // 6:00 PM
10 }

```

h31/ x +

```

~/workspace/solutions/h31/ $ clang++ -c tester.cpp
tester.cpp:9:7: error: 'hours' is a private member of 'Time'
    t.hours = 18;    // 6:00 PM

```

By prohibiting clients from directly accessing **private** data, the implementation can assume that all access to **private** data goes through the **public** interface (unlike the **Time struct** of last week, where clients **should use the member functions**, but **were not prohibited** from directly accessing the data members **hours** and **minutes**.)

Actually, the only difference between **class** and **struct** in C++ is that with a **struct**, the members are **public** by default; with a **class** they are **private**. By convention, we will use **struct** for **POD** (plain-old-data) data types, and **class** for encapsulated types.

The Implicit Parameter

Consider the **getHours()** member function of the **Time** class:

```
int Time::getHours() const
{
    return m_hours;
}

int main()
{
    Time t;                                // a Time object
    cout << t.getHours() << endl;         // value of t::m_hours
}
```

The **getHours()** member function **does not** contain a local variable named **m_hours**. But, the function still compiles and runs correctly. Why?

In a **member function**, you may **directly access and manipulate** any or all of the class's **data members** by referring to them by name. You don't indicate that **m_hours** is a data member, nor do you specify **which Time** object you're referring to.

C++ assumes that all data members **are the data members of the receiver object**, and so the line **return hours** means "return the value of the **hours** data member of the object on which this function was invoked." In such a case, **the receiver object is known as the implicit parameter**, passed to every member function.

The Pointer **this**

Behind the scenes, the implicit parameter is a **pointer to the calling object**. Every member function has an implicit parameter. Thus the effective signature for the **getHours()** function is as if you had declared it like this:

```
int getHours(const Time* const this);
```

The keyword **this** is the **automatically supplied name** for the implicit parameter. The **const** following the **Time*** means that the value inside the pointer can never be changed; it always points to the block of data containing the object's

data members. The **const** following the member function header means that the implicit parameter is a pointer to a **const Time** object.

If you wish, you can **explicitly** use the pointer when calling other member functions, or accessing data members:

```
int Time::getHours() const
{
    return this->m_hours;
};
```

How this is Initialized

When you **call a member function** like this:

```
Time t;                                // a Time object
cout << t.getHours() << endl;         // value of t::m_hours
```

That call is **implicitly translated** into code that acts as if you had written:

```
Time t;                                // a Time object
cout << getHours(&t) << endl;         // value of t::m_hours
```

Because of this call, the **this** pointer is initialized to the **address of the calling object**.

The sum() Member Function

Consider the **sum()** member function from **Time**:

```
class Time
{
public:
    Time sum(const Time& rhs) const;
    . . .
};
```

When you add two **Time** objects (**a + b**) together like this:

```
Time after = a.sum(b);
```

The **caller** (the implicit parameter) is the left-hand-side of the expression **a + b**.

Thus, the effective implicit prototype for the function is similar to this:

```
Time sum(const Time* lhs, const Time& rhs);
```

In the implementation, however, instead of the `explicit lhs` parameter shown here, you'd use the keyword `this`.

```
Time Time::sum(const Time& rhs) const
{
    auto tMinutes = this->m_hours * 60 + this->m_minutes;
    auto dMinutes = rhs.m_hours * 60 + rhs.m_minutes;
    . . .
}
```

If you leave off the keyword `this`, it is assumed. Notice that when you implement a `const` member function, you `repeat` the word `const` in the implementation.

Mutators & Constructors

Many classes have `set*` member functions. These are called **mutators**, since they **change the state** of the object. Mutators should **validate data** written to the object to **enforce the class invariants**. With properly written mutators, the errors described earlier **cannot occur**. Consider your `Time` class. If you add `setHours()` and `setMinutes()` members to the class, you would have to enforce these restrictions:

- **hours** must be between **0** and **23** inclusive
- **minutes** and **seconds** must both be between **0** and **59** inclusive

Unlike the `read()` member function, where you could put the stream into a failed state, if these conditions were not met, in a mutator you need to **throw** an exception like this:

```
void Time::setHours(int h)
{
    if (h < 0 || h > 23) throw out_of_range("...");
    m_hours = h;
}
```

Getter & Setter Patterns

The pattern of pairing a **get*** along with a **set*** function is common. It is always safer to allow clients to **read the values** of data members than it is allow them to **change** those values. Thus, setter methods are less common than getters in class design. Classes with no mutators at all, are called **immutable classes**.

Unlike Java, the actual **get*** and **set*** name pattern is not as common. Instead, what programmers often do is write a **pair of overloaded functions**. The accessor is **const** and returns a value. The non-**const** mutator **const** returns a reference, which can be assigned to. Instead of the name **getHours()**, use the name **hours()** for both of them.

```
Time Time::hours() const { return m_hours; } // accessor
Time& Time::hours()           // mutator
{
    if (h < 0 || h > 23) throw out_of_range("...");
    return m_hours;
}
```

Constructors

Initializing object data is the responsibility of the **constructor**, which always has the same name as the class and **never has a return type**. A constructor is a member function which **initializes an object into a well-formed state** before clients start manipulating it. When C++ creates an object from a class:

1. It **allocates** a block of memory large enough to store the data elements
2. It passes **the address of that block** of memory to the constructor function.
The address is the **this** pointer inside the constructor function.

The constructor **is called automatically** whenever an object is created. If you have a class that defines a constructor, that constructor is **guaranteed to execute** whenever you create an object of the class type.

Default Constructors

The **default constructor** is the constructor which takes **no arguments** and which should initialize **all of its data members** to an appropriate **default** value. Alternatively, since C++11, you **provide an initial value** when defining the data members, just as in Java.

If you do not provide a constructor, the compiler will "write" one for you. This is called the **synthesized default constructor**. If you use in-class initializers, then this is perfect.

Member Initialization

In C++, all constructors must initialize **all primitive types**. A constructor does not need to initialize any object members (like string or vector). This is **exactly the opposite from Java**, where you must initialize all of the object instance variables, or they are set to null (an invalid object). Primitive instance variables are automatically initialized to **0**.

```
public class Point
{
    private String name;
    int x, y;
    public Point() {}
}

Point p = new Point(); // x,y->0, name is null (invalid)
```

In C++, if you fail to initialize a primitive data member, then it assumes **whatever random value** was in memory; if you don't initialize an object, such as **string** or **vector**, its default constructor will run, and it is still a valid object.

```
class Point
{
    string name;
    int x, y;
    Point() {}
};

Point p; // x,y->random, name is empty string
```

Of course, if you provide default initializers for your primitive data members, they **will** automatically be initialized.

Working Constructor

With the **Time** we might like to have another, **overloaded** constructor which takes hours, minutes and seconds. This is generally known as the **working constructor**. Here is the public interface of **Time** with both of these constructors.

```
class Time
{
public:
    // Constructors
    Time(); // default
    Time(int h, int m); // working
    // . . .
};
```

Unfortunately, if you have **any** explicit constructors, the synthesized one **is deleted**, so you have to add an **explicit default constructor**. In C++11, however, you can just add the phrase **=default**; to the end of the prototype in the class header, and the compiler will **retain** the synthesized constructor that normally writes.

Implementing Constructors

The **implementation** of the constructors goes into the **.cpp** file along with the other member functions. The job of the constructor is to **initialize the data members**, so in the `Time` class, you might have code that looks something like this.

```
Time::Time() { m_hours = m_minutes = 0; }
Time::Time(int h, int m)
{
    m_hours = h; m_minutes = m;
}
```

The Constructor Initializer List

Look at these two statements:

```
string a = "Bob", b;    // initialize
b = "Bill";             // assign
```

Two **string** objects are created and initialized on line one; **a** is initialized using the C-String **"Bob"**, and **b** is initialized to the empty **string** by running the default constructor.

On line 2, the **string** object **b** is destroyed, a new **string** object is initialized with **"Bill"**, and that new **string** object replaces the **string** object originally held by **b**. The variable **b** is first initialized, then destroyed, then assigned.

The body of the constructor is executed **after** the memory for the data members have already been allocated. You may use **assignment** to place a new value into these data members, just like you would in Java.

For primitive types, the cost of doing this is negligible, but for object types, such assignments mean that **data members are constructed twice**—once at allocation and once at initialization. Here's an example. (The implementation is inline to shorten the code.)

```
class Person
{
    string name;
public:
    Person(const string& n) { name = n; }
};
```


When you write `Person p("Fred")`, the `name` data member first calls the default constructor to create an empty string object. Then, in the body of the constructor, the default-constructed `string` is destroyed when assigning `n` to `name`. This is inefficient, and you want to avoid it.

We can instruct the compiler to **initialize the individual data members** at the time that the data members are **allocated** instead. This is called the **initializer list**:

- It follows the parameter list and is preceded by a colon (:)
- It is followed by a list of member names and their initializers.
- Initialization occurs **in the order the members are declared in the class**.

In C++98 the initializers are placed in parentheses; in C++11 use either parentheses or braces. You cannot use the assignment operator. Here is the same class using the initializer list. In this case, the `name` data member is **only constructed once**:

```
class Person
{
    string name;
public:
    Person(const string& n) : name(n) {}
};
```

Delegating Constructors

C++11 added a feature called **delegating constructors**, which allows a constructor to use another constructor for its actual work. Usually the working constructor is the one that is implemented, with the others simply **forwarding** their requests. The purpose of this is to **eliminate redundant code** in multiple constructors.

For instance, consider the working constructor from the `Time` class.

```
Time::Time(int h, int m)
{
    m_hours = h; m_minutes = m;
}
```

Using delegation, the default constructor can be written like this:

```
Time::Time() : Time(0, 0)
{ }
```

Finish Up

- Complete the **reading exercises (REX)** for this chapter.
- Complete the homework using the **CS50 IDE**. The link is on Canvas.
 - a. Make sure you **submit** the assignment using **make submit**.
 - b. Make sure you check the [CS150 Homework Console](#) to see that your scores got reported, **before** the beginning of the next lecture.
- Take the **pre-class reading quiz** on Canvas. You have two attempts.

See you in class or on the Canvas discussion board.