

Polymorphism

Value-Oriented or **Object-Based** programming involves creating new, user-defined types. There are three strategies for building a new type:


- Build it **completely from scratch**, using only built-in components.
- **Combine simpler types** to create complex types. This is **composition**.
- **Extend** a general class, adding new features. This is called **inheritance**.

Programming with inheritance is called **Object Oriented Programming**.

Has-A and Is-A

These strategies express two kinds of "class relationships":

- The **has-a relationship**, says one class is a combination of other objects. In the **has-a** relationship one type is composed of different parts. A **Bicycle** class thus may contain two instances of the **Wheel** class.
- The **is-a relationship**, when one class is an extension or "kind of" another class. The **is-a** relationship occurs when members of one class are a **subset** of another class. This is also called an **inheritance** relationship. In the inheritance relationship shown here, we'd say that a **MountainBike** is-a **Bicycle**.



Polymorphic Inheritance

Inheritance is a form of **specialization**. The derived class inherits both the member functions and the data members from the base class, while optionally adding more of both. The derived class **IS-A specialized form** of the more general base class.

A derived class **may override** a **virtual** member function to add specialized behavior, as we did with **Student::toString()**. This is called **polymorphic inheritance**, it provides specialized behavior in response to the same messages. Let's see if that's true. Let's use our simple **Person←Student** hierarchy from the last chapter and see what happens with some experiments.

A Student *is-a* Person





Click the Running Man on the left to open a copy of the lab. Change `toString()` in each class so it identifies the class at the beginning of the method. Here are the modified `toString()` member functions. Notice that this version of the `Student::toString()` no longer calls its base class version; it **entirely replaces it**. I've highlighted the changes.

```
string Person::toString() const
{
    return "Person::Name: " + name;
}

string Student::toString() const
{
    return "Student::Name: " + getName()
        + ", ID: " + to_string(studentID);
}
```

Use **make run** to see the client program running. This is a kind of "polymorphism", known as static polymorphism. You send the same message to different objects and each **responds to the same message** according to its nature.



This is not what we mean when we talk about **polymorphism**. This works just as if **Person** and **Student** were not related in any way.

Polymorphism means an inheritance relationship where the request can be sent to any kind of **Person** object, and the specialized **Persons**, such as **Students** or **Employees** respond appropriately.

The Slicing Problem

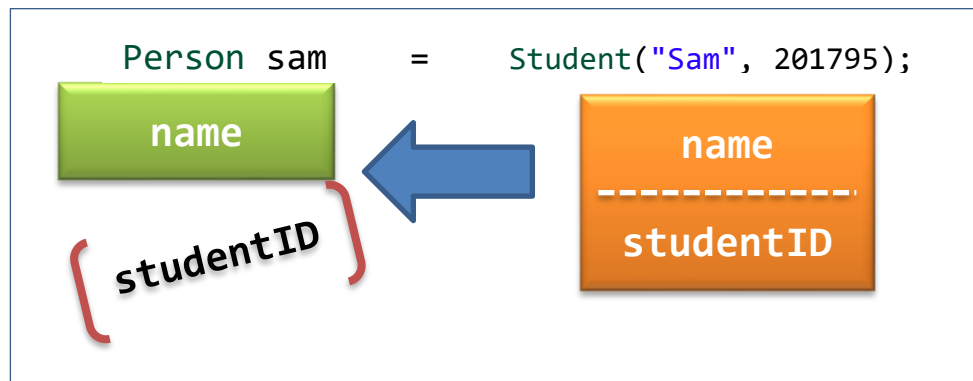
Change the example (`client.cpp`) again, so it looks like this:

```
int main()
{
    Person sam = Student("Sam", 201795);
    Person pam = Person("Pam B.");
    ...
}
```

Now you have two **Person** objects: one "plain" **Person**, a one specialized **Person** who is a **Student**. Is the output the same as previously? **Of course not!**

For the **Student sam**, you know longer see the **ID**. And, both the **Student** and the **Person** are identified with **Person::Name**, even though we do have an overridden member function, **Student::toString()**. Here's why this happens.

First, objects in C++ are **value types**, unlike the reference types in Java. When you assign a derived class object to a base class variable, **only the base class portion** of the object is copied. **This is called the slicing problem.**




If you pass a derived class object **by value** to a function that expects a base class object, the same slicing will occur as well. This is easy to fix. Just **always** follow this rule:

NEVER EVER EVER EVER ASSIGN A DERIVED CLASS OBJECT TO A BASE CLASS VARIABLE. EVER!

No Polymorphic Objects

While slicing is a problem it is not the real culprit. The problem is **polymorphism only works in C++ with references or pointers**. Fix the **client.cpp** like this:

```
Student sam = Student("Sam", 201795);
Person pam = Person("Pam B.");
Person& samRef = sam;
Person& pamRef = pam;
cout << "sam says->" << samRef.toString() << endl;
cout << "pam says->" << pamRef.toString() << endl;
```



Now, the **Person&** reference **samRef** refers to the **Student** object **sam**, and when we call **samRef.toString()** it calls **Student::toString()**, not **Person::toString()** like our previous examples did.

Polymorphic Functions

What we really want are **polymorphic functions** like this:

```
// A polymorphic function
void greet(const Person& p)    // any kind of Person
{
    cout << "Hello, I'm " << p.toString() << endl;
}
```

This function is polymorphic because the formal parameter is a **reference to a base class**. (Note, **not** a base class object.) You can **pass any kind of Person**, such as **Student** or **Employee** objects and they will behave appropriately.

Polymorphic functions should operate on references or pointers to a base class. Functions should never use pass-by-value with base class objects.

Polymorphic Lists

Creating a list (**vector** or array) of different kinds of object also leads to slicing:

```
vector<Person> v;
v.push_back(Student("Sam", 201795));
v.push_back(Person("Pam B."));
```

When you **push_back** the **Student** or **Employee** objects, **they are sliced** as well. The **vector v** **does not** contain a **Student** and a **Person**; it contains two **Person** objects. Sam has been stripped of everything that makes him a **Student**; he has been **effectively lobotomized**; he no longer knows who he is.

No Lists of References

You also cannot fall back on using references, like you did with polymorphic functions, since you cannot create a **vector<Person&> v** or an array **Person& a[3]**. **Both of these declarations are illegal**. A reference is not a variable or object (**Lvalue**), but an **alias** for an existing **Lvalue**.



Pointers to the Rescue?

One solution is to create a `vector<Person*> v` or array `Person * a[2]`. Here's a short example that places two different kinds of `Person` pointers in a `vector` and prints them. Each person responds appropriately. Go ahead and add this code to `main()`. Include the `<vector>` header.

```
int main()
{
    vector<Person*> people;
    people.push_back(new Student("Sam", 201795));
    people.push_back(new Person("Pam B.));

    for (auto p : people) {cout << p->toString() << endl;}
    for (auto p : people) delete p;
}
```

Since two of these objects are created on the heap, it is up to you to reclaim their memory before the `vector` goes out of scope and it is lost. The `vector` cannot do it because it does not know if the pointers it contains point to objects on the heap or objects on the stack. If you add a stack-based pointer to this program, it crashes.

Smart Pointers to the Rescue

You can eliminate the need to reclaim memory by using the two C++11 smart pointers, `shared_ptr` and `unique_ptr`, which are declared in the header `<memory>`. Here is a version that uses `unique_ptr`, which doesn't require the `delete` loop at the end.

```
using up = unique_ptr<Person>;
vector<up> people;
people.push_back(up{new Student("Sam", 201795)});
people.push_back(up{new Person("Pam B.)});
for (auto& p : people) cout << p->toString() << endl;
```

The `unique_ptr` constructor is `explicit`, so a `type alias` helps to shorten the name. Also, `unique_ptr` objects cannot be copied, so the range-based `for` loop accesses each element by reference, which is what you need anyway.


Because the `vector` now stores **only smart pointers**, you can't inadvertently add a stack-based pointer to the list.

Inheritance Forms

What would happen if you were to remove the keyword **virtual** from the definition of the **toString()** member function in the **Person** class? Your code would still compile, but the **toString()** function would no longer be overridden; it **would be hidden** in the derived class **Student**.


Functions are **bound** to an object depending on how they are declared. A non-virtual function is **bound at compile time** to the class that it is defined in. A non-virtual function defined in the **Person** class (such as **getName()**), will always be bound to the **Person** class, and cannot be overridden in any subsequent classes.

This is called **early binding** (or compile-time binding).



When you add the keyword **virtual** to a function, the function call is not determined (or bound) at compile time, but **when the program is run**. Instead of looking at the type of the pointer or reference used in the function call, **the actual object pointed to** is used to decide which function to call. This decision is made when your program runs. If your **Shape*** points to a **Circle** object, then **Circle::draw()** will be called, but **only if draw()** is a **virtual** function.

This is called **late-binding** or **dynamic dispatch**.¹ In Java, **all** methods use late binding, but in C++ you, as the base-class designer get to decide which version to use, through the application of the keyword **virtual**.



Virtual member functions are implemented by adding a new pointer to every object that contains at least one **virtual** function. This pointer is called a **vp**tr and it points to a table of functions, called a **vtable**.² The *vtable* contains the actual addresses of the functions to be called for that class.

¹ Wikipedia – Dynamic Dispatch: http://en.wikipedia.org/wiki/Dynamic_dispatch

² Wikipedia – Virtual Method Table http://en.wikipedia.org/wiki/Virtual_method_table

Using this illustration, let's see how late binding, or dynamic dispatch works:

1. You call `emp[0]->raiseSalary()`
2. Your call is routed through the `vptr` in `emp[0]`, which is actually a `manger`, and eventually finds the address of the `Manager::raiseSalary()` function inside the `Manager` vtable.
3. You call `emp[1]->promote()`
4. Your call is routed through the `vptr` in `emp[1]`, which is actually an `Engineer`. This `vptr` points to the `Engineer` vtable where it finds the `Engineer::promote()` method.

Multiple Inheritance

C++ includes a capability known as **multiple inheritance**, which allows a class to be derived from two (or more) base classes. Multiple inheritance lets you create a class **AmphibiousVehicle** from the parents **LandVehicle** and **WaterVehicle**.




An **AmphibiousVehicle** inherits wheels or treads from its **LandVehicle** parent, and a prop or screw and rudder from its **WaterVehicle** parent. You can't do this in Java, because Java classes can have only one parent class.

MI in the Stream Libraries


The standard stream libraries include classes that are **both** input and output streams.

The **fstream** class at the bottom is an **iostream**, which is in turn—if you follow the arrow leading up and to the left—an **istream**. The **fstream** class therefore inherits all the methods that pertain to **istreams**. If you instead follow the arrow up and to the right, you discover that the **fstream** class is an **ostream**, which means that it inherits these methods as well.



Implementation Inheritance

Suppose you have a class which simulates a **Cadillac**. It has an exceptionally fine sound system, which required a lot of effort to implement and of which you're especially proud. Now you want a **GiPod** class, for sale over the Web:



Because you've already created the **Cadillac** class, why not just create a derived class, and then eliminate all the member functions that have nothing to do with playing music, **transforming the car into a mere sound system**?

To reuse the code you've already written, you replace **brake()**, **accelerate()**, and all the other "extra" methods from the **Cadillac** class with empty braces. In traditional computer-science terms, you replace them with a **NOP** (No Operation).

This practice, called **contraction**, is a trap! You should avoid doing this for two reasons:

- You're **violating the substitutability rule**. You will undoubtedly break some code that relied on all **Cadillac** objects carrying out certain operations.
- It's **more work than doing the right thing!**

Private Inheritance

Private inheritance is one way to solve this problem. Private inheritance means you want to **inherit the implementation of a class, not the interface**. A class has some functionality that you want to exploit, but you don't want to use the interface of the base class.

```
class GiPod : private Cadillac {};
```


GiPod objects would not, in the **is-a** sense, be **Cadillac** objects. Calling any "inherited" member functions would fail. To call any of the inherited methods, you must add those methods to the new interface, with a **using declaration** like this:

```
class GiPod : private Cadillac
{
public:
    using Cadillac::playMusic;
};
```

You don't need to specify the arguments or supply an argument list. You do need the **public:** if you want the name moved into the **public** section.

Doing this, you have reduced the interface to a single method. The relationship between the **Cadillac** and **GiPod** classes is one of "**implemented in terms of**", not "**is-a**".

Composition Revisited

Perhaps a better solution is to use **composition**. **Composition creates a new class by combining simpler classes**, using instances of the simpler class as the data members. In composition, an object is **composed of other objects**, which make up its "parts." That's why it's called (informally) a **has-a** relationship; because we can say that:

- A car **has a** motor, or
- A bicycle **has a** seat, or
- A computer **has a** CPU



Here is version of the **GiPod** that uses composition; note that it needs to **explicitly write the prototypes** for any methods that it uses.

```
class GiPod
{
public:
    void playMusic() const
    {
        caddy.playMusic(); // forward or delegate
    }
private:
    Cadillac caddy;        // data member
};
```

Specification Inheritance



A final form of inheritance is called **specification inheritance**. A base class may **specify a set of responsibilities** that a subclass **must** fulfill, but not provide any actual implementation. The **interface** (method signatures) are inherited.

The specification relationship is used **in combination** with regular specialization:


- the derived class **inherits the interface** of the base class, as in specification
- it also inherits a **default implementation** of, at least, some of its methods

A derived class **may** **override** a **virtual** member function to add specialized behavior, as we did with **Student::toString()**, or, it may be **required** to implement a particular member function, which could not be provided in the base class.

Abstract Classes

The classes we've used so far are called **concrete** classes. We can also create **abstract classes**. An abstract class is usually an **incomplete** class, a class that contains certain methods that are **specified**, but not **implemented** in the definition of the class.

Because the abstract class contains these incomplete methods, **it cannot be used to create objects**—it can only be used as a base class when defining other classes. That is, it **only** makes sense in the context of **polymorphic inheritance**.




Suppose you have an **abstract Shape** base class that doesn't have the faintest notion of how to implement its abstract **draw()** method, yet it knows that each of its **concrete** derived classes will need to do so.

Only the **concrete** classes derived from **Shape**—**Circle**, **Square**, and **Star**—possess the necessary knowledge to actually **draw()** themselves. You, only need to program in terms of **Shape** objects; the actual shapes will take care of their own behavior.

Pure Virtual Functions

In C++, an **Abstract Base Class** (or **ABC**) is any class that has one, or more, **pure virtual member functions**, created using the following syntax in the prototype:

```
class Shape    // abstract class
{
public:
    ...
    // Pure virtual function (abstract method)
    virtual void draw() const = 0;
    ...
};
```




Think of the **= 0;** part of syntax as a **replacement** for the **abstract** keyword in Java. Abstract classes are **not restricted** to abstract member functions like **draw()**; you can have as many regular (concrete) member functions as you'd like, freely mixed with your abstract methods. The **Shape** class in the UML diagram at the left has a **setLocation()** member function. In UML, abstract methods, such as **draw()**, are drawn using italics.

Your concrete methods may **call** abstract methods as part of their definition, even though the member function is never implemented in the base class. C++ pure virtual functions **may have an optional implementation**. Since you cannot create an instance of an abstract base class, you must call this implementation from a derived class.

Using an Abstract Class

It is **illegal to create an instance of an abstract class**. Your compiler enforces this. You may, however, create **ABC** pointers or references as long as they point to, or refer to concrete objects which are derived from the **ABC**.



When you **extend an abstract class**, your derived class **must override each and every abstract function in its base class**, giving each a concrete implementation. The resulting derived class is a **concrete class**, and it can be used to create new objects.

Abstract classes thus provide a way of **guaranteeing** that an object of a given type will understand a given message. In that sense, **they specify a set of responsibilities that a derived class must fulfill**.

A Triangle Example

Let's look at an example. To create the **Triangle** (or **Circle** or **Square**) classes, using the abstract **Shape** class as the base class, all you need do is:

```
class Triangle : public Shape
{
public:
    // MUST override; pure virtual in Shape class
    void draw() const;
};

void Triangle::draw() const { /* your code here */ }
```

1. Specify the **Shape** class as the **public** base class in the class header.
2. Provide an implementation for **every** abstract method in the **Shape** class.

For **Triangle** that means you **must** define a **draw()** member function where indicated by the comments. In reality, you'll probably do a lot more; **Circle** might have a **radius** member, **Square** class could have members for the **size** of each side, and the **Triangle** class could have members for **base** and **height**.

Final Members

A derived class may **replace** an inherited member function, which the class designer wanted left alone. In the **Person** class, **getName()** returns the name of the person. **This works fine as is**; there's no reason for a derived class to change it.

However, **nothing prevents a derived class from redefining it** like this:

```
class Imposter : public Person
{
public:
    string getName() const
    {
        return "Emperor " + Person::getName();
    }
};
```

The derived class has no access to the **private** member **name**. However, because the **getName()** member function can be **redefined**, the derived class was able to effectively gain access to this field. And, because of the **principle of substitutability**, the **Person** that your function receives as a parameter may actually be an **Imposter**.

To prevent this, when you design a base class, consider which functions you want to allow others to extend and which ones should be "set in stone". No one should ever change **getName()**, so you can **seal it using final** like this:

```
class Person
{
public:
    virtual std::string getName() const final;
};
```

When a member function is marked **final** then derived classes are **prevented** from overriding it and we would see an error message like this:

```
Person.cpp:33:12: error: virtual function 'virtual
std::string Imposter::getName() const'
string getName() const
      ^
Person.cpp:21:8: error: overriding final function
'virtual std::string Person::getName() const'
string Person::getName() const { return name; }
      ^
```

Final Classes

Only **virtual** functions can be marked **final**. When designing a collection of classes, you normally **won't want all of the classes to be extensible**. To prevent others from extending your class, add **final** to the class header, just like you did for the method.

If you make a class **final**, then there is no reason to make the methods **final** as well.

Finish Up

- Complete the **reading exercises (REX)** for this chapter.
- Complete the homework using the **CS50 IDE**. The link is on Canvas.
 - a. Make sure you **submit** the assignment using **make submit**.
 - b. Make sure you check the [CS150 Homework Console](#) to see that your scores got reported, **before** the beginning of the next lecture.
- Take the **pre-class reading quiz** on Canvas. You have two attempts.

See you in class or on the Canvas discussion board.